

▼ Connecting with google drive

```
from google.colab import drive
drive.mount('/content/drive/')

Mounted at /content/drive/

datadir = '/content/drive/My Drive/CS444/assignment4_materials/assignment4_materials/'
import os
if os.path.exists(datadir):
    os.chdir(datadir)
else:
    print(datadir, 'path not found')

!pwd
/content/drive/My Drive/CS444/assignment4_materials/assignment4_materials

# run this cell only once
# !unzip cats.zip
```

▼ Generative Adversarial Networks

For this part of the assignment you implement two different types of generative adversarial networks. We will train the networks on a dataset of cat face images.

```
import torch
from torch.utils.data import DataLoader
from torchvision import transforms
from torchvision.datasets import ImageFolder
import matplotlib.pyplot as plt
import torch.nn as nn
import numpy as np

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

%load_ext autoreload
%autoreload 2

from gan.train import train

device = torch.device("cuda:0" if torch.cuda.is_available() else "gpu")
```

▼ GAN loss functions

In this assignment you will implement two different types of GAN cost functions. You will first implement the loss from the [original GAN paper](#). You will also implement the loss from [LS-GAN](#).

▼ GAN loss

TODO: Implement the `discriminator_loss` and `generator_loss` functions in `gan/losses.py`.

The generator loss is given by:

$$\ell_G = -\mathbb{E}_{z \sim p(z)} [\log D(G(z))]$$

and the discriminator loss is:

$$\ell_D = -\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] - \mathbb{E}_{z \sim p(z)} [\log(1 - D(G(z)))]$$

Note that these are negated from the equations presented earlier as we will be *minimizing* these losses.

HINTS: You should use the `torch.nn.functional.binary_cross_entropy_with_logits` function to compute the binary cross entropy loss since it is more numerically stable than using a softmax followed by BCE loss. The BCE loss is needed to compute the log probability of the true label given the logits output from the discriminator. Given a score $s \in \mathbb{R}$ and a label $y \in \{0, 1\}$, the binary cross entropy loss is

$$bce(s, y) = -y * \log(s) - (1 - y) * \log(1 - s)$$

Instead of computing the expectation of $\log D(G(z))$, $\log D(x)$ and $\log(1 - D(G(z)))$, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing.

```
from gan.losses import discriminator_loss, generator_loss
```

Least Squares GAN loss

TODO: Implement the `ls_discriminator_loss` and `ls_generator_loss` functions in `gan/losses.py`.

We'll now look at [Least Squares GAN](#), a newer, more stable alternative to the original GAN loss function. For this part, all we have to do is change the loss function and retrain the model. We'll implement equation (9) in the paper, with the generator loss:

$$\ell_G = \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)) - 1)^2]$$

and the discriminator loss:

$$\ell_D = \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} [(D(x) - 1)^2] + \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)))^2]$$

HINTS: Instead of computing the expectation, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing. When plugging in for $D(x)$ and $D(G(z))$ use the direct output from the discriminator (`scores_real` and `scores_fake`).

```
from gan.losses import ls_discriminator_loss, ls_generator_loss
```

GAN model architecture

TODO: Implement the Discriminator and Generator networks in `gan/models.py`.

We recommend the following architectures which are inspired by [DCGAN](#):

Discriminator:

- convolutional layer with `in_channels=3, out_channels=128, kernel=4, stride=2`
- convolutional layer with `in_channels=128, out_channels=256, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=256, out_channels=512, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=512, out_channels=1024, kernel=4, stride=2`
- batch norm
- convolutional layer with `in_channels=1024, out_channels=1, kernel=4, stride=1`

Use padding = 1 (not 0) for all the convolutional layers.

Instead of Relu we LeakyReLu throughout the discriminator (we use a negative slope value of 0.2). You can use simply use relu as well.

The output of your discriminator should be a single value score corresponding to each input sample. See `torch.nn.LeakyReLU`.

Generator:

Note: In the generator, you will need to use transposed convolution (sometimes known as fractionally-strided convolution or deconvolution).

This function is implemented in pytorch as `torch.nn.ConvTranspose2d`.

- transpose convolution with `in_channels=NOISE_DIM, out_channels=1024, kernel=4, stride=1`
- batch norm
- transpose convolution with `in_channels=1024, out_channels=512, kernel=4, stride=2`
- batch norm
- transpose convolution with `in_channels=512, out_channels=256, kernel=4, stride=2`
- batch norm
- transpose convolution with `in_channels=256, out_channels=128, kernel=4, stride=2`
- batch norm
- transpose convolution with `in_channels=128, out_channels=3, kernel=4, stride=2`

The output of the final layer of the generator network should have a `tanh` nonlinearity to output values between -1 and 1. The output should be a $3 \times 64 \times 64$ tensor for each sample (equal dimensions to the images from the dataset).

```
from gan.models import Discriminator, Generator
```

Data loading

The cat images we provide are RGB images with a resolution of 64x64. In order to prevent our discriminator from overfitting, we will need to perform some data augmentation.

TODO: Implement data augmentation by adding new transforms to the cell below. At the minimum, you should have a RandomCrop and a ColorJitter, but we encourage you to experiment with different augmentations to see how the performance of the GAN changes. See <https://pytorch.org/vision/stable/transforms.html>.

```
batch_size = 32
imsize = 64
cat_root = './cats'

cat_train = ImageFolder(root=cat_root, transform=transforms.Compose([
    transforms.ToTensor(),

    # Example use of RandomCrop:
    transforms.Resize(int(1.15 * imsize)),
    transforms.RandomCrop(imsize),
]))

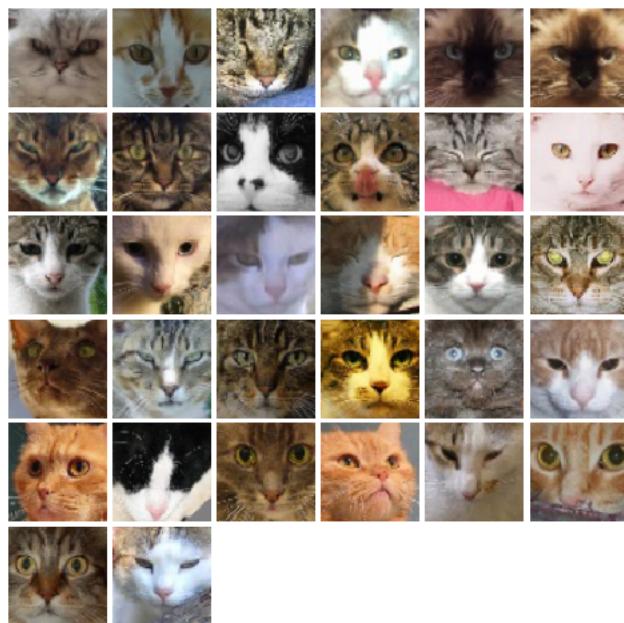
cat_loader_train = DataLoader(cat_train, batch_size=batch_size, drop_last=True)
```

▼ Visualize dataset

```
from gan.utils import show_images
```

```
imgs = next(cat_loader_train.__iter__())[0].numpy().squeeze()
show_images(imgs, color=True)

/usr/local/lib/python3.9/dist-packages/torchvision/transforms/functional.py:1603:
    warnings.warn('
```



▼ Training

TODO: Fill in the training loop in `gan/train.py`.

```
NOISE_DIM = 100
NUM_EPOCHS = 50
learning_rate = 1e-3
```

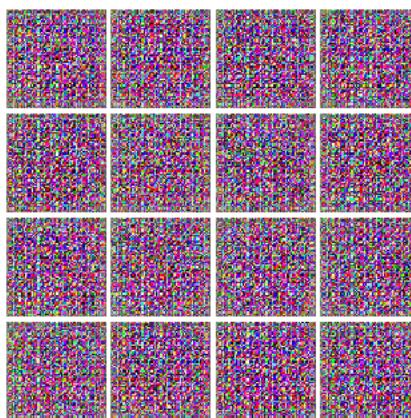
▼ Train GAN

```
D = Discriminator().to(device)
G = Generator(noise_dim=NOISE_DIM).to(device)

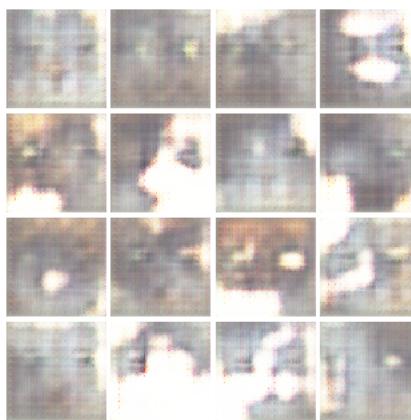
D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))
```

```
# original gan
train(D, G, D_optimizer, G_optimizer, discriminator_loss,
      generator_loss, num_epochs=NUM_EPOCHS, show_every=250,
      batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

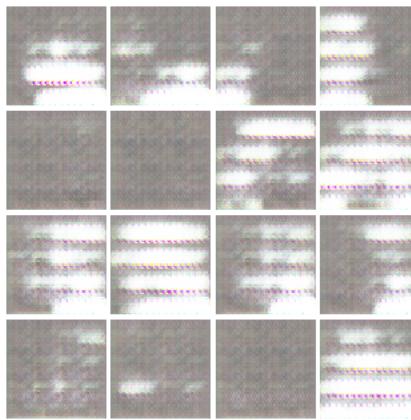
EPOCH: 1
/usr/local/lib/python3.9/dist-packages/torchvision/transforms/functional.py:1603: UserWarning: The default value of the antialias parameter has changed from None to 'bicubic'.
warnings.warn(
Iter: 0, D: 1.404, G:3.864



Iter: 250, D: 1.82, G:0.5296



EPOCH: 2
Iter: 500, D: 0.9001, G:1.798



Iter: 750, D: 1.015, G:1.361

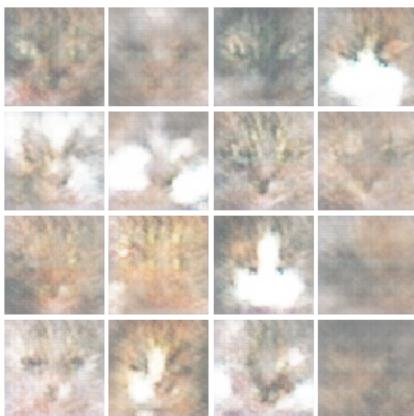


EPOCH: 3
Iter: 1000, D: 0.7664, G:3.154





Iter: 1250, D: 0.8841, G:1.086



EPOCH: 4

Iter: 1500, D: 0.8442, G:2.22



Iter: 1750, D: 0.8579, G:1.805



EPOCH: 5

Iter: 2000, D: 0.8601, G:1.994





Iter: 2250, D: 0.349, G:2.903



EPOCH: 6
Iter: 2500, D: 1.063, G:3.722



Iter: 2750, D: 1.334, G:4.081



EPOCH: 7
Iter: 3000, D: 0.7145, G:3.047





Iter: 3250, D: 0.5861, G:3.654



EPOCH: 8

Iter: 3500, D: 0.6522, G:2.149



Iter: 3750, D: 0.2513, G:2.74



EPOCH: 9

Iter: 4000, D: 0.4427, G:2.123



Iter: 4250, D: 0.3313, G:2.422



EPOCH: 10

Iter: 4500, D: 0.2112, G:4.82



Iter: 4750, D: 0.4704, G:5.033



EPOCH: 11

Iter: 5000, D: 0.3216, G:6.3



Iter: 5250, D: 0.2058, G:4.732





EPOCH: 12
Iter: 5500, D: 0.3856, G:7.39



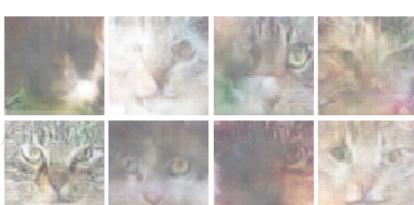
Iter: 5750, D: 0.1506, G:3.392



EPOCH: 13
Iter: 6000, D: 0.2137, G:4.73



Iter: 6250, D: 0.08691, G:5.7





EPOCH: 14
Iter: 6500, D: 0.04791, G:5.405



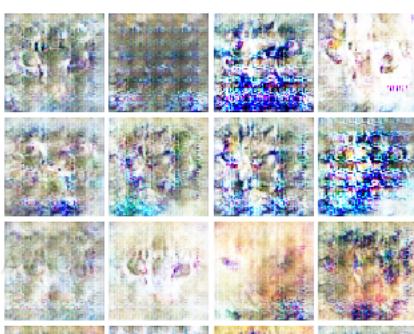
Iter: 6750, D: 0.7524, G:9.153



EPOCH: 15
Iter: 7000, D: 0.02584, G:5.369

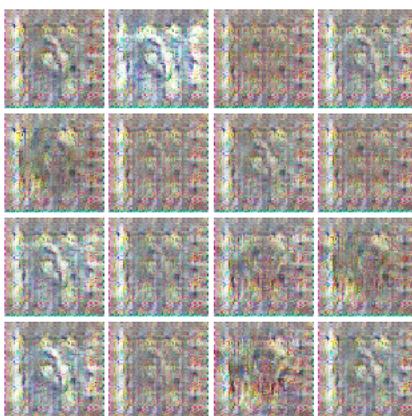


Iter: 7250, D: 0.02313, G:6.916

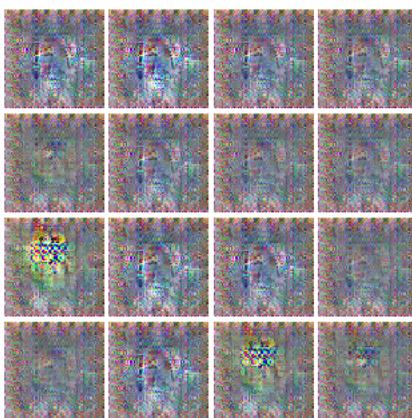




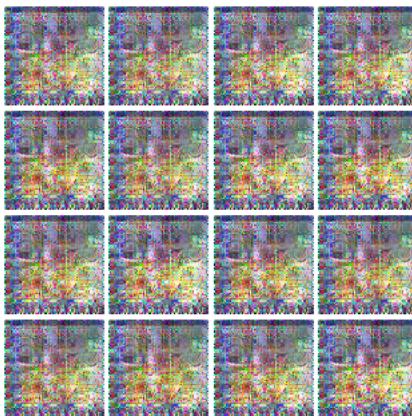
EPOCH: 16
Iter: 7500, D: 0.007139, G:7.551



Iter: 7750, D: 0.0005296, G:8.18



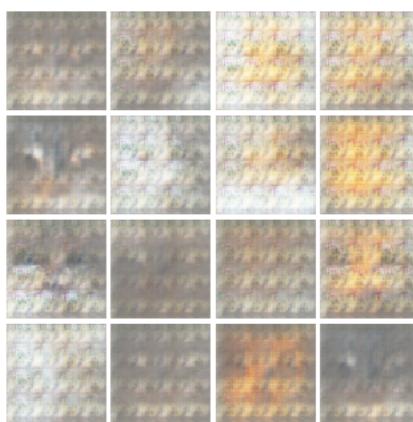
EPOCH: 17
Iter: 8000, D: 1.976e-05, G:15.48



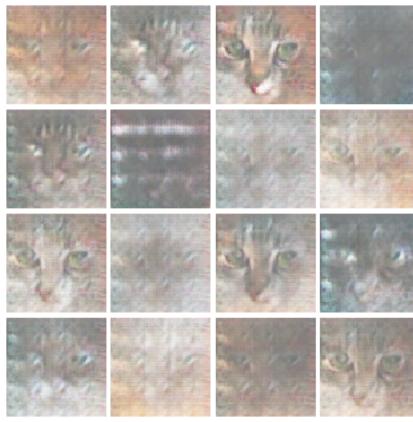
Iter: 8250, D: 0.0001648, G:39.14



EPOCH: 18
Iter: 8500, D: 0.1004, G:6.548



Iter: 8750, D: 0.03448, G:7.636



EPOCH: 19
Iter: 9000, D: 0.466, G:7.725



Iter: 9250, D: 0.1047, G:8.315



EPOCH: 20
Iter: 9500, D: 0.06076, G:6.453





Iter: 9750, D: 0.1272, G:10.18



EPOCH: 21

Iter: 10000, D: 0.02186, G:8.294

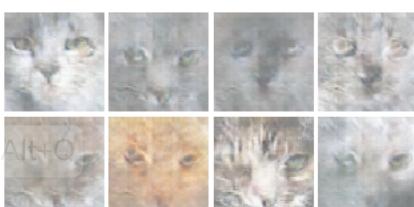


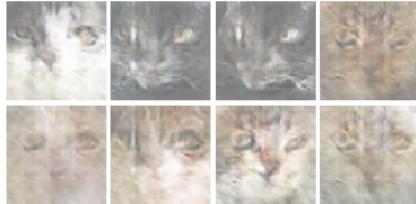
Iter: 10250, D: 0.4661, G:7.568



EPOCH: 22

Iter: 10500, D: 0.03727, G:5.514





Iter: 10750, D: 3.177, G:23.89



EPOCH: 23

Iter: 11000, D: 0.0986, G:8.857



Iter: 11250, D: 0.0135, G:10.74



EPOCH: 24

Iter: 11500, D: 0.1785, G:6.648





Iter: 11750, D: 0.2865, G:11.29



EPOCH: 25

Iter: 12000, D: 0.05508, G:8.072



Iter: 12250, D: 0.02403, G:8.981



EPOCH: 26

Iter: 12500, D: 0.08988, G:5.659



Iter: 12750, D: 2.436, G:11.41



EPOCH: 27
Iter: 13000, D: 0.08051, G:8.076



Iter: 13250, D: 0.09018, G:5.813



EPOCH: 28
Iter: 13500, D: 0.4847, G:9.353



Iter: 13750, D: 0.1462, G:8.17





EPOCH: 29
Iter: 14000, D: 0.1872, G:10.31



Iter: 14250, D: 0.04759, G:7.341



EPOCH: 30
Iter: 14500, D: 0.03389, G:6.373



Iter: 14750, D: 0.2461, G:9.222





EPOCH: 31
Iter: 15000, D: 0.05453, G:6.652



Iter: 15250, D: 0.04684, G:11.24



EPOCH: 32
Iter: 15500, D: 0.1645, G:8.869



EPOCH: 33
Iter: 15750, D: 0.02804, G:8.139



Iter: 16000, D: 0.03255, G:8.812



EPOCH: 34

Iter: 16250, D: 0.0445, G:8.608



Iter: 16500, D: 0.03176, G:9.594



EPOCH: 35

Iter: 16750, D: 0.01037, G:6.301



AHQ

Iter: 17000, D: 0.003622, G:9.171



EPOCH: 36
Iter: 17250, D: 6.267, G:22.55



Iter: 17500, D: 0.027, G:6.338



EPOCH: 37
Iter: 17750, D: 0.2946, G:10.88



Iter: 18000, D: 0.02904, G:6.504





EPOCH: 38
Iter: 18250, D: 0.1327, G:5.132



Iter: 18500, D: 0.04387, G:7.312



EPOCH: 39
Iter: 18750, D: 0.01074, G:5.997



Iter: 19000, D: 0.1913, G:9.741





EPOCH: 40
Iter: 19250, D: 0.004206, G:7.149



Iter: 19500, D: 0.9517, G:17.4



EPOCH: 41
Iter: 19750, D: 0.01317, G:8.667



Iter: 20000, D: 0.2456, G:11.78



Alt+Q

EPOCH: 42
Iter: 20250, D: 0.1777, G:6.409



Iter: 20500, D: 0.3138, G:9.189



EPOCH: 43
Iter: 20750, D: 0.3407, G:8.388



Iter: 21000, D: 0.1599, G:8.553



EPOCH: 44
Iter: 21250, D: 0.1078, G:4.975





Iter: 21500, D: 0.1653, G:9.83



EPOCH: 45
Iter: 21750, D: 0.002121, G:6.347



Iter: 22000, D: 0.1816, G:12.32



EPOCH: 46
Iter: 22250, D: 0.01723, G:7.671





Iter: 22500, D: 0.02963, G:19.37



EPOCH: 47

Iter: 22750, D: 0.005308, G:11.33



Iter: 23000, D: 0.08054, G:10.02



EPOCH: 48

Iter: 23250, D: 0.05585, G:9.492





Iter: 23500, D: 0.005678, G:5.476



EPOCH: 49

Iter: 23750, D: 0.01358, G:9.314



▼ Train LS-GAN



```
D = Discriminator().to(device)
G = Generator(noise_dim=NOISE_DIM).to(device)

D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.999))
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.999))
```

```
# default betas = 0.5, 0.999
```

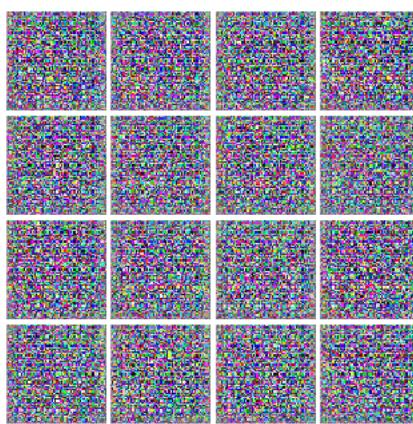


```
# ls-gan
train(D, G, D_optimizer, G_optimizer, ls_discriminator_loss,
      ls_generator_loss, num_epochs=NUM_EPOCHS, show_every=250,
      batch_size=batch_size, train_loader=cat_loader_train, device=device)
```

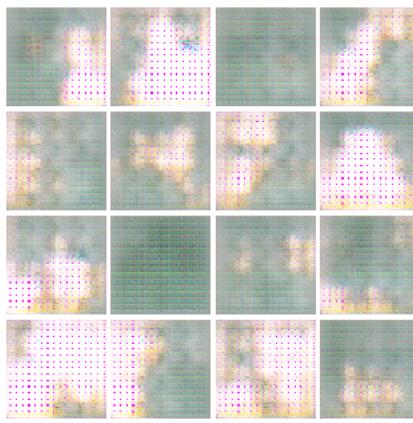
▷

Alt+Q

EPOCH: 1
Iter: 0, D: 0.9051, G:31.82



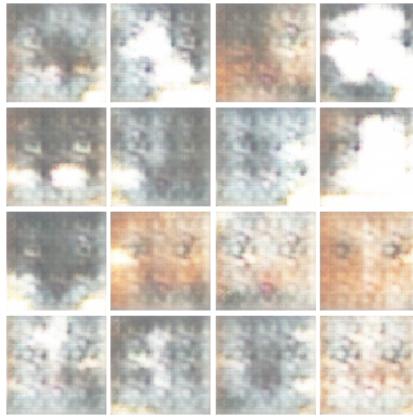
Iter: 250, D: 0.494, G:0.6876



EPOCH: 2
Iter: 500, D: 0.4639, G:0.2442



Iter: 750, D: 0.3506, G:0.6265

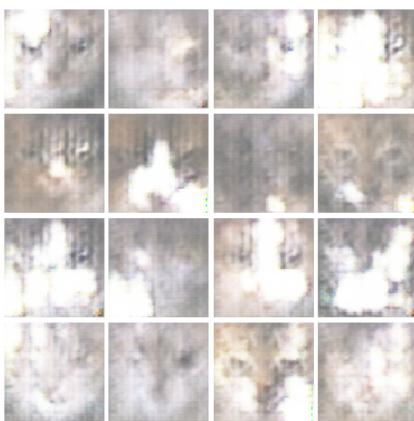


EPOCH: 3
Iter: 1000, D: 0.2665, G:0.3978

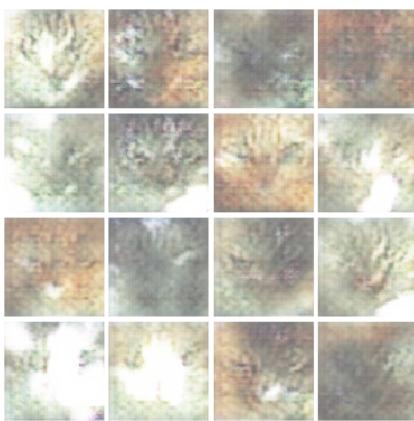




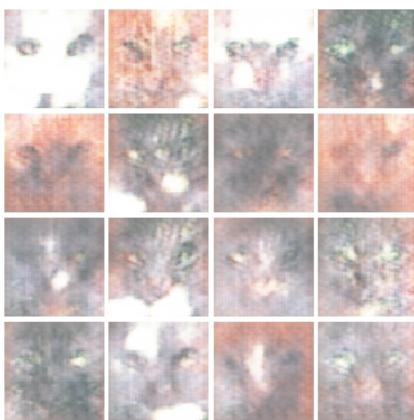
Iter: 1250, D: 0.3603, G:0.4334



EPOCH: 4
Iter: 1500, D: 0.3042, G:1.345



Iter: 1750, D: 0.8207, G:0.1738



EPOCH: 5
Iter: 2000, D: 0.5936, G:0.4628





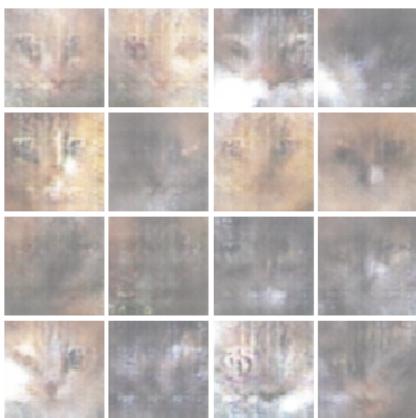
Iter: 2250, D: 0.3505, G:0.4596



EPOCH: 6
Iter: 2500, D: 0.3537, G:0.7555



Iter: 2750, D: 0.1911, G:1.006



EPOCH: 7
Iter: 3000, D: 0.2527, G:0.6386





Iter: 3250, D: 0.2866, G:0.6723



EPOCH: 8

Iter: 3500, D: 0.36, G:0.445



Iter: 3750, D: 0.2878, G:0.525



EPOCH: 9

Iter: 4000, D: 0.1532, G:0.5567



Iter: 4250, D: 0.00432, G:0.8286

Iter: 4250, D: 0.09452, G:0.8260



EPOCH: 10
Iter: 4500, D: 0.1983, G:1.146



Iter: 4750, D: 0.254, G:0.6202



EPOCH: 11
Iter: 5000, D: 0.07278, G:1.055



Iter: 5250, D: 0.2014, G:0.7008





EPOCH: 12
Iter: 5500, D: 0.1222, G:0.8312



Iter: 5750, D: 0.08405, G:0.9092



EPOCH: 13
Iter: 6000, D: 0.3684, G:0.5358



Iter: 6250, D: 0.1844, G:1.079





EPOCH: 14
Iter: 6500, D: 0.1286, G:0.802



Iter: 6750, D: 0.1936, G:0.8203



EPOCH: 15
Iter: 7000, D: 0.2907, G:1.259



Iter: 7250, D: 0.159, G:0.6063





EPOCH: 16
Iter: 7500, D: 0.169, G:1.102



Iter: 7750, D: 0.1915, G:0.7845



EPOCH: 17
Iter: 8000, D: 0.1903, G:1.657



Iter: 8250, D: 0.1658, G:1.045



Alt+Q

EPOCH: 18
Iter: 8500, D: 0.1273, G:0.9983



Iter: 8750, D: 0.1297, G: 0.9305



EPOCH: 19

Iter: 9000, D: 0.09855, G: 0.7985



Iter: 9250, D: 0.04403, G: 0.8487



EPOCH: 20

Iter: 9500, D: 0.1306, G: 0.7928





Iter: 9750, D: 0.1545, G:1.442



EPOCH: 21
Iter: 10000, D: 0.07784, G:1.216



Iter: 10250, D: 0.06578, G:0.7414



EPOCH: 22
Iter: 10500, D: 0.09487, G:0.831





Iter: 10750, D: 0.059, G:1.202



EPOCH: 23

Iter: 11000, D: 0.1423, G:0.8791



Iter: 11250, D: 0.1439, G:1.763



EPOCH: 24

Iter: 11500, D: 0.1507, G:0.9113



Iter: 11750, D: 0.107, G:0.6437



EPOCH: 25

Iter: 12000, D: 0.1261, G:1.438



Iter: 12250, D: 0.06329, G:0.9867



EPOCH: 26

Iter: 12500, D: 0.04847, G:0.9894



Iter: 12750, D: 0.07956, G:1.033





EPOCH: 27
Iter: 13000, D: 0.03513, G: 1.349



Iter: 13250, D: 0.09473, G: 0.6015



EPOCH: 28
Iter: 13500, D: 0.1331, G: 0.5



Iter: 13750, D: 0.04539, G: 0.7968





EPOCH: 29
Iter: 14000, D: 0.1089, G:0.7395



Iter: 14250, D: 0.1057, G:0.8082



EPOCH: 30
Iter: 14500, D: 0.05947, G:1.212



Iter: 14750, D: 0.0768, G:1.13



AIoQ



EPOCH: 31
Iter: 15000, D: 0.09253, G:1.106



Iter: 15250, D: 0.03511, G:1.012



EPOCH: 32
Iter: 15500, D: 0.1081, G:1.48



EPOCH: 33
Iter: 15750, D: 0.05206, G:1.446



Iter: 16000, D: 0.1769, G:1.498



EPOCH: 34

Iter: 16250, D: 0.03723, G:1.288



Iter: 16500, D: 0.0335, G:0.813



EPOCH: 35

Iter: 16750, D: 0.05509, G:0.8695

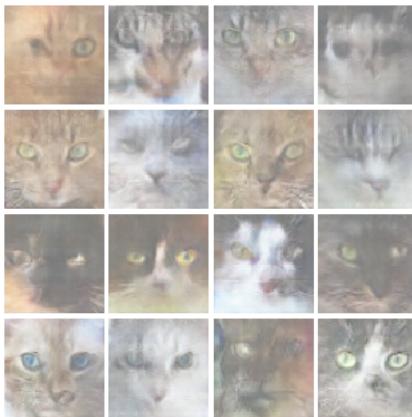


Iter: 17000, D: 0.1509, G:0.8049





EPOCH: 36
Iter: 17250, D: 0.1419, G:1.125



Iter: 17500, D: 0.08103, G:0.7166

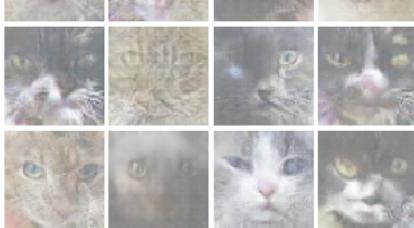


EPOCH: 37
Iter: 17750, D: 0.07883, G:1.496



Iter: 18000, D: 0.08668, G:1.157





EPOCH: 38
Iter: 18250, D: 0.151, G:0.5944



Iter: 18500, D: 0.0334, G:1.16



EPOCH: 39
Iter: 18750, D: 0.04385, G:1.244



Iter: 19000, D: 0.03921, G:0.9263





EPOCH: 40
Iter: 19250, D: 0.09095, G:1.682



Iter: 19500, D: 0.04769, G:0.8927



EPOCH: 41
Iter: 19750, D: 0.06901, G:0.6549



Iter: 20000, D: 0.2912, G:0.6517



EPOCH: 42
Iter: 20250, D: 0.04952, G:0.9108



Iter: 20500, D: 0.06106, G:1.261



EPOCH: 43
Iter: 20750, D: 0.04882, G:1.073



Iter: 21000, D: 0.03985, G:0.8783



EPOCH: 44
Iter: 21250, D: 0.06278, G:0.8178





Iter: 21500, D: 0.05479, G: 0.8499



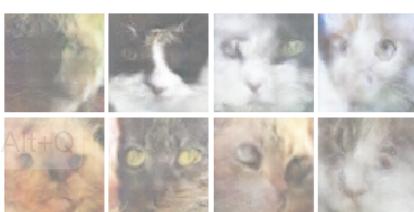
EPOCH: 45
Iter: 21750, D: 0.03921, G: 0.8772



Iter: 22000, D: 0.09923, G: 0.8288



EPOCH: 46
Iter: 22250, D: 0.07719, G: 0.6311





Iter: 22500, D: 0.09058, G:1.017



EPOCH: 47

Iter: 22750, D: 0.0205, G:1.088



Iter: 23000, D: 0.03951, G:1.003



EPOCH: 48

Iter: 23250, D: 0.06864, G:0.9358



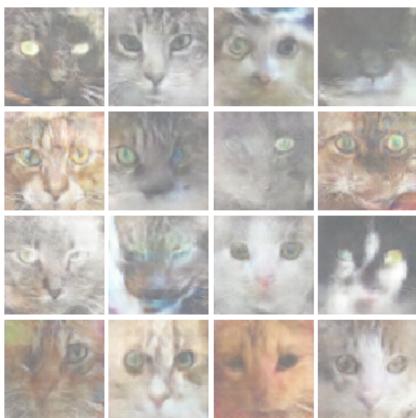


Iter: 23500, D: 0.2035, G:0.5662



EPOCH: 49

Iter: 23750, D: 0.03067, G:0.9909



Iter: 24000, D: 0.09764, G:0.6439



EPOCH: 50

Iter: 24250, D: 0.03836, G:0.908



Iter: 24500 D: 0.1092 G:0.7194

18077 24566, D: 0.1052, S: 0.7151



[Colab paid products - Cancel contracts here](#)

